

SHAILESH DUDALA

SENIOR APPLIED AI ENGINEER

shaileshdudala@icloud.com · github.com/smgpulse007 · linkedin.com/in/ssdudala

I build governed intelligent systems that move from uncertain signals to typed, validated workflows with explicit human authority. My work spans predictive healthcare, local document intelligence, context-engineered claims agents, and evidence-gated open-source execution harnesses.

EXPERIENCE

MetLife via Bizintex — Applied AI Engineer Consultant

2026 — present

Designing public-safe, context-engineered document and claims workflows with typed state, tool contracts, validation, exception routing, reviewer authority, and trace telemetry.

- ≈90% lower document-handling effort in a measured claim-packet workflow.
- ≈50% shorter time-to-claim-payable in a supported workstream.

Inland Empire Health Plan via Infowave — Lead DS/ML Engineer Consultant

2023 — 2025

Led local RAG/OCR, healthcare quality evidence extraction, predictive ML, FWA analytics, MLOps, and operational reporting inside controlled data boundaries.

- Cleared a 7K-case review backlog with ≈90% lower review time.
- Improved automated measure closures by 20%; reduced transportation waste by 18%.

Hexplora — Lead Data Scientist / Product

2020 — 2023

Built and scaled a predictive analytics platform across nine healthcare programs, combining calibrated risk models, data products, and care-manager workflows.

- Supported \$500K in new revenue and ≈\$3M in aggregate client performance-based payouts.
- Reduced model-deployment cycles by ≈50%.

CommonSpirit Health · Health New England · University of Chicago · AbbVie — Data science and biomedical informatics

Earlier foundations

Worked across hospital analytics, provider data, clinical sensors, genomics, biomedical research, and public-health modeling.

- Analytics spanning 142 hospitals; 100K+ provider records validated.
- Reduced manual provider-review effort by 60%.

OPEN-SOURCE SYSTEMS

META HARNESS	LLM STEERING
Pre-1.0 evidence-gated execution harness for coding-agent phases, bounded packets, validation, proof ledgers, checkpoints, and continuation. Source-installed and read-only-first.	Activation-steering workbench with model registry, prompt-pair controls, pre/post hooks, layer/coefficient sweeps, and validated result artifacts.

CORE SYSTEMS

Context-engineered agents · Predictive ML and calibration · Local OCR/RAG and evidence lineage · Healthcare interoperability (FHIR/HL7) · Typed validation and human review · LLMOps/MLOps · Evaluation, replay, golden sets, shadow/canary release · Cost, latency, tool, and failure telemetry

EDUCATION + CREDENTIALS

M.S., Biomedical Informatics — University of Chicago, 2019 · B.Tech., Computer Science & Engineering — SRM University, 2018 · Yale Summer School, Public Health Modeling, 2019

AWS Certified Machine Learning Engineer — Associate · Microsoft Azure AI Fundamentals (AI-900) · Microsoft Power BI Data Analyst (PL-300)

PUBLIC EDITION · No phone, precise location, private application data, confidential payloads, employer prompts, PHI, PII, or production secrets.